

Session 2

AI and Machine Learning

Hult International Business School

Michael de la Maza

Version 1.1




Advanced Decision Trees: Agenda



Regularization



Boosting



Bagging



Application
using XGBoost

What is a decision tree?

- A tree-like structure in which:
 - Internal nodes represent tests on an attribute
 - Each branch represents the outcome of the test
 - Each leaf node represents a class label (in the case of a classifier) or a number (in the case of a regressor)



Regularization



Regularization intuition: Reduce the risk of overfitting by preferring 'small' models

- Two general approaches:
 - Penalize 'large' models
Example: In linear regression, prefer models with fewer input variables.
 - Use hyperparameters to 'force' models to be small
Example: Limit the depth of decision trees



Regularization hyperparameters in decision trees

max_depth:

- If set, then the maximum depth of the tree cannot exceed this parameter.
- If not set, then tree continues to split until nodes are all of the same class (or until the number of samples is less than min_samples_split)

min_samples_split:

- The minimum number of samples needed for a node to be split. If a node has fewer than this number of examples, it will not be split.
- The default value is 2. So, by default, it will continue to split until there is only one sample.


max_leaf_nodes:

Sets the maximum number of leaf nodes.

min_impurity_decrease:

A node will not be split if the purity decrease is less than this value.

There are several other regularization hyperparameters.



Note: Make sure to add hyperparameter to your glossary

Let's look at code

- Optimize `max_depth` (just like last session)
- Optimize `min_weight_fraction_leaf` (require a minimum number of samples at each leaf)
- Optimize `min_impurity_decrease` (require impurity to decline by a minimum amount)

End result is a smaller tree with same performance as original tree



5 minute exercise

- Select another regularization hyperparameter and train a decision tree with a specific value for that hyperparameter
- Examples:
 - `min_samples_split = 20`
 - `min_samples_leaf = 10`
 - `max_leaf_nodes = 6`



5 minute exercise

- Optimize the hyperparameter
(e.g., check values between 1 and 10)



Boosting



Boosting Intuition

- Train predictors (classifiers/regressors) in sequence with each subsequent predictor attempting to correct errors made by the previous predictors.
- Predictors are combined by weighing predictions according to accuracy: the higher the accuracy, the higher the weight the predictor receives.
- Example:
 - Train decision tree A on data.
 - Train decision tree B on data and increase the weight of instances that A gets wrong.



AdaBoost Classifier

- AdaBoostClassifier


`DecisionTreeClassifier(max_depth=1, n_estimators=5, learning_rate=0.5, random_state=42)`




Note: It is normal to have 50+ estimators. We use 5 so we can visualize.

5 minute exercise: Boosting

- Try AdaBoost with different $n_estimators = 2, 4, 6, 8$
 - How does accuracy change?
 - Which one would you choose?
- Advanced: Write a for loop over $n_estimators$ from 1 to 10
 - How do you plot the decision trees?
 - What is the best value of $n_estimator$?



Bagging




Intuition

- Train same classifier on different subsets of the training set.
 - Sample instances
 - Sample features
- Produce estimate by aggregating predictions:
 - For classifiers, choose the most popular class (or aggregate probabilities).
 - For regressors, average the predictions.



5 minute exercise

- What do you expect the result to be when:
 - $n_estimators = 1$
 - $max_samples = 142$



XGBoost



XGBoost: 'extreme' decision trees

- Combines many good ideas for creating decision trees. Focuses on boosting, as the name implies.
- Was the core algorithm in many winning machine learning contests.
- XGBoost is the default model for structured classification problems.



Decision Tree Review




What is a decision tree?

- A. A method for optimizing model parameters
- B. A tree-like structure used to make predictions based on input features
- C. A technique for finding the optimal solution in a search space
- D. A tool for data visualization




What is the primary advantage of using decision trees in machine learning?

- A. They are easy to interpret and visualize
- B. They can handle missing data
- C. They have high predictive accuracy



What are the two types of nodes in decision trees?

- A. Red nodes and blue nodes
- B. Round nodes and square nodes
- C. Decision nodes and leaf nodes



What does this graph indicate?

- A. Overfitting
- B. Underfitting
- C. The training error decreases as the validation error increases
- D. The training error increases as the validation error decreases

